

Bubble + IQN

Ttobot

{ sejunpark, dujung, victor.cappa } @kaist.ac.kr

1. Base code

- dopamine by google
<https://github.com/google/dopamine>
- use basic code + config (DQN, Rainbow, IQN)
- customized Retro Environment and IQN Agent



2. Reward

- bonus on score and killing enemy. but penalty on death

```
def _calculate_step_reward(self, curr_level, curr_score, curr_lives, curr_enems, game_over):  
    """  
    reward for score configuration:  
    [objective]  
    - survive as long as possible  
    - achieve as much as score  
    """  
    # init with base penalty per each step  
    acc_rew = self.step_penalty  
    # kill an enemy - with double reward along with score (1 kill -> 100 score)  
    if self.last_enems > curr_enems:  
        acc_rew += 1 * (self.last_enems - curr_enems)  
    # case of life lost  
    if self.last_lives > curr_lives:  
        acc_rew += -5 * (self.last_lives - curr_lives)    # max 3x lives  
    # get enhancement in log scale.  
    if curr_score > (self.last_score + 1):  
        acc_rew += math.log(curr_score - self.last_score, 100) + self.score_bonus  
    # return with total.  
    return acc_rew
```

step-penalty:
0.0005

score_bonus:
0.02

3. Action

- Discrete action-space with key combinations

```
# NOTE - core actions for BubbleBobble.  
self.mapping = {  
    0: [0, 0, 0, 0, 0, 0, 0, 1, 0], # RIGHT  
    1: [0, 0, 0, 0, 0, 0, 1, 0, 0], # LEFT  
    2: [0, 0, 0, 0, 0, 0, 0, 0, 1], # JUMP  
    3: [1, 0, 0, 0, 0, 0, 0, 0, 0], # FIRE  
    4: [1, 0, 0, 0, 0, 0, 0, 1, 0], # FIRE + RIGHT  
    5: [1, 0, 0, 0, 0, 0, 1, 0, 0], # FIRE + LEFT  
}
```

4. Pre-processing

- Green channel as grey scaled image by 84x84

```
def _fetch_grayscale_observation(self, obs, output):  
    # clear walls  
    for wall in self.last_walls:  
        # masked = np.all(obs == wall, axis=-1)  
        # obs[masked] = [255,32,32]  
        obs[np.all(obs == wall, axis=2)] = [255,62,62]  
    # use Green channel as grayscale (SIMPLE BUT FAST)  
    obs = obs[:, :, 1]  
    np.copyto(output, obs)  
    return output
```

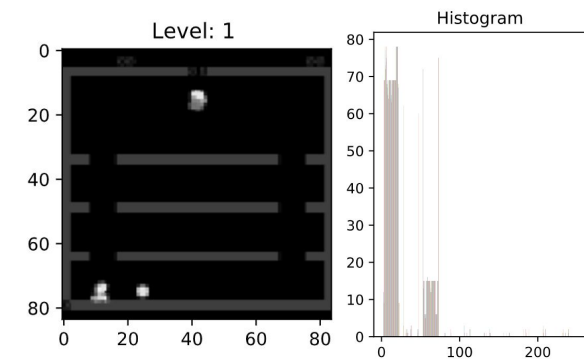

You, a month ago • chore: format indent

```
def _pool_and_resize(self):  
    # Pool if there are enough screens to do so.  
    if self.frame_skip > 1:  
        np.maximum(  
            self.screen_buffer[0], self.screen_buffer[1], out=self.screen_buffer[0])  
  
    transformed_image = cv2.resize(self.screen_buffer[0],  
                                   (self.screen_size, self.screen_size),  
                                   interpolation=cv2.INTER_AREA)  
    int_image = np.asarray(transformed_image, dtype=np.uint8)  
    return np.expand_dims(int_image, axis=2)
```

5. Performance


- GTX 1080 Ti (12GB) + 24xCore CPU w/ 64GB RAM
- 250k steps / train. 10k steps / eval.

	steps/sec	iteration/day	rewards/train
No Wall Filter	90	~25	-
Wall Filter	60 (-30%)	~17	10% improved



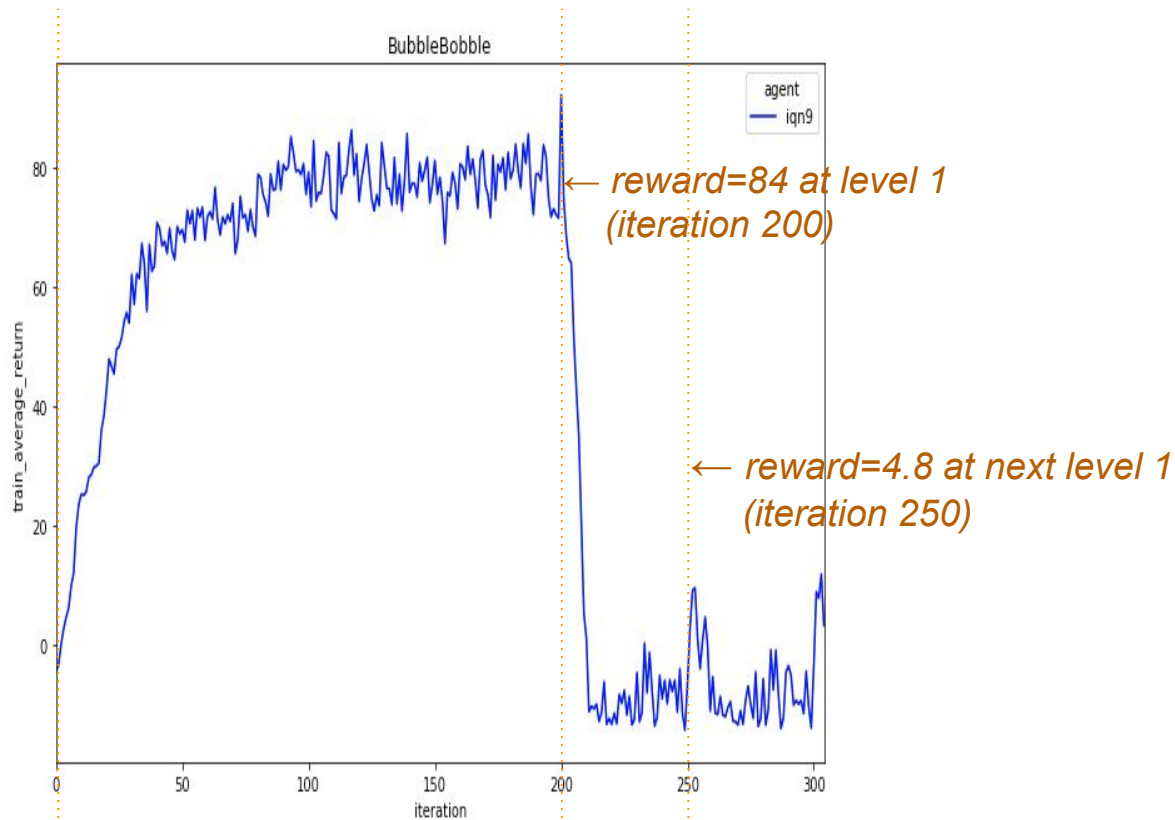
7. Experimental

- Prepared the trained IQN-Agent with 1 level by 200 iteration.



8. Train Forgot?

- Reward at Level 1. Down to 4.8 from 84 after training each 50 levels in sequence



APPENDIX

Implicit Quantile Network (IQN) Tensorflow Agent playing BubbleBobble game

